

# Order-k Markov Transformer LPE Experiments (k=1,2,3)

Automated run via lpe/markov\_k\_transformer.py

February 15, 2026

## Summary

This report summarizes order-k experiments for  $k \in \{1, \dots, 7\}$  with independent Beta(1,1) priors per transition state. Training used sequence length equal to rollout length:  $100 \cdot 2^k$ .

## Core Metrics

| $k$ | Model NLL | Bayes NLL | Gap (%) | Posterior MAE | LPE Median             | RelErr (%) |
|-----|-----------|-----------|---------|---------------|------------------------|------------|
| 1   | 0.491818  | 0.487222  | 0.94    | 0.0160        |                        | 100.00     |
| 2   | 0.504101  | 0.493396  | 2.17    | 0.1642        | 2970363860660420608.00 |            |
| 3   | 0.498595  | 0.489566  | 1.84    | 0.0647        |                        | 128.16     |

Table 1: Primary metrics by order  $k$ .

## Figures

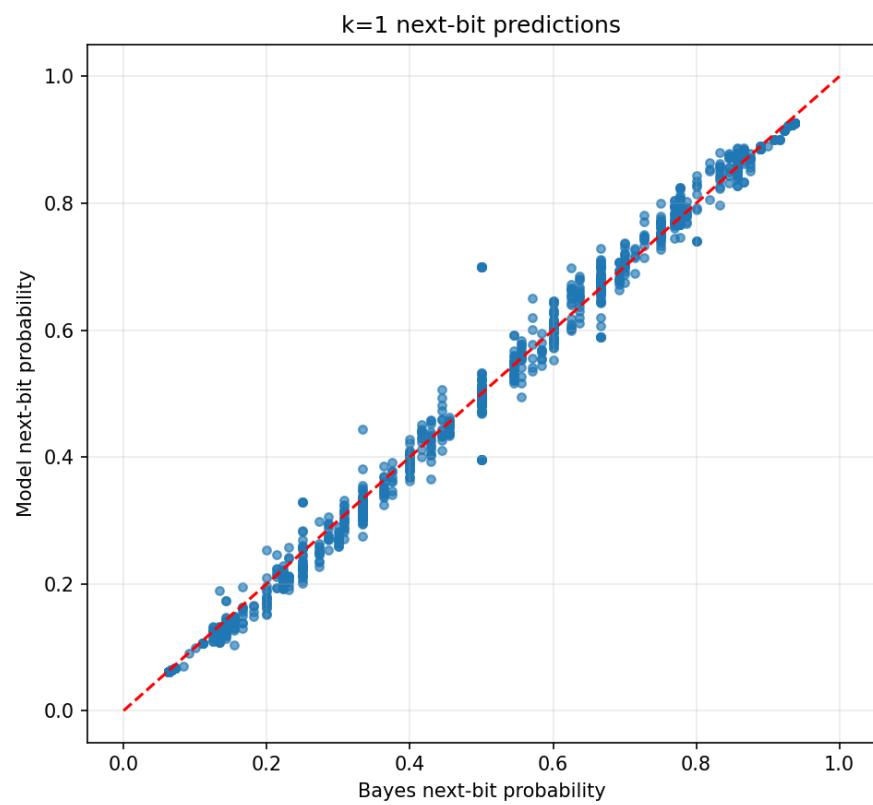


Figure 1:  $k=1$ : Step 1: next-bit prediction vs Bayes

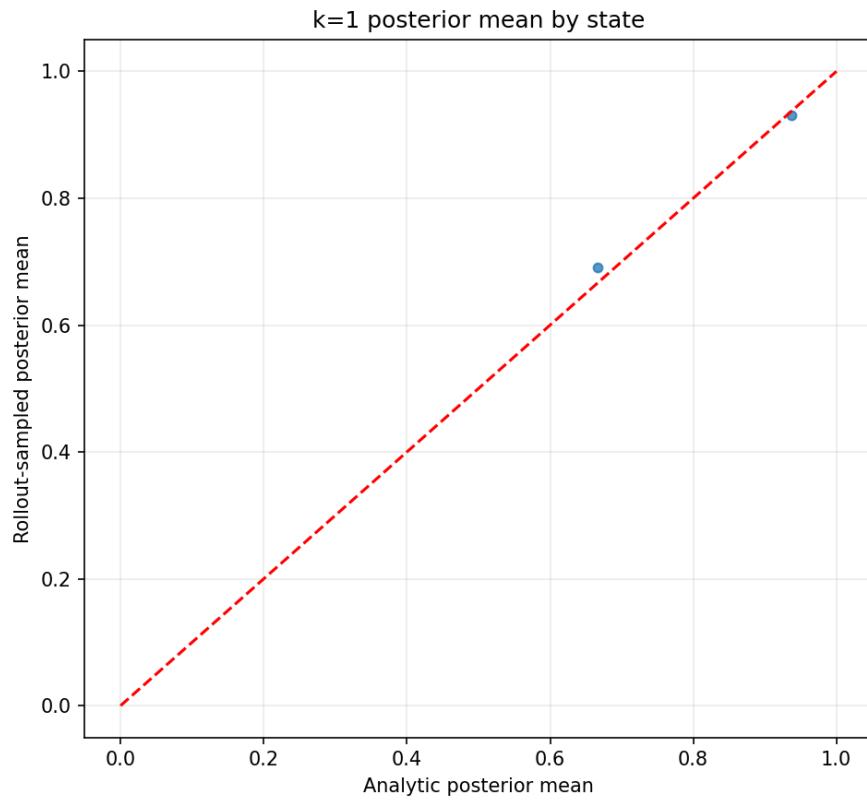


Figure 2:  $k=1$ : Step 2: posterior mean comparison

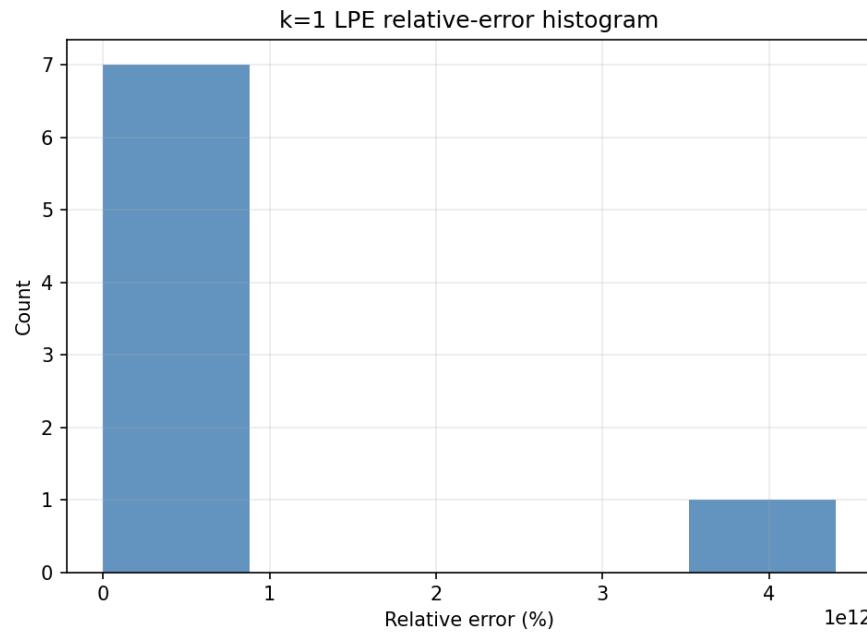


Figure 3:  $k=1$ : Step 3: LPE relative-error histogram

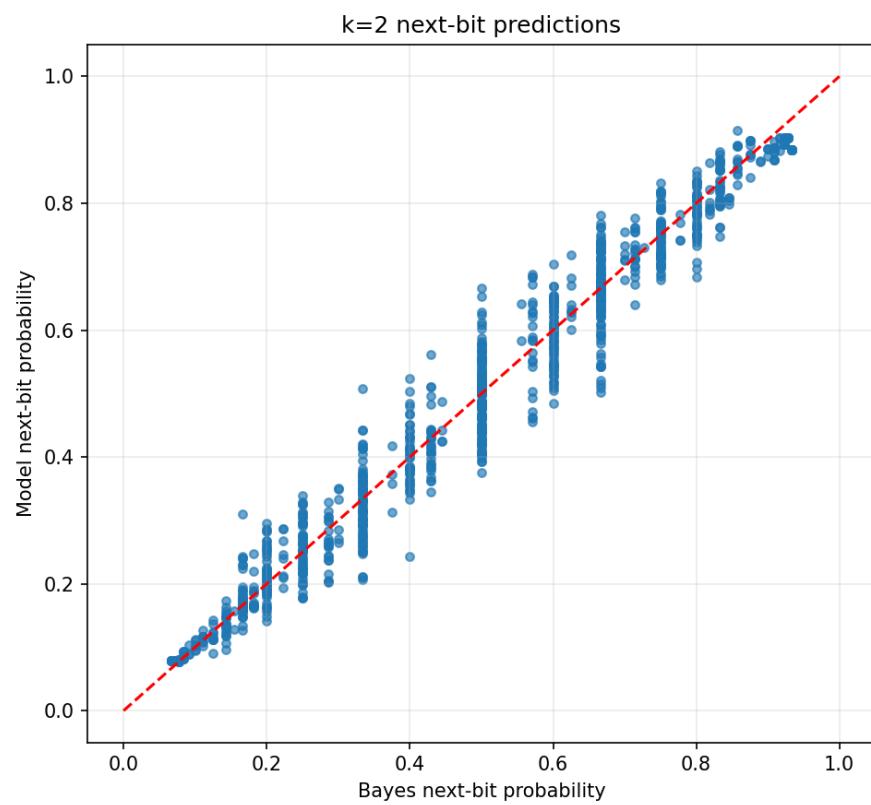


Figure 4: k=2: Step 1: next-bit prediction vs Bayes

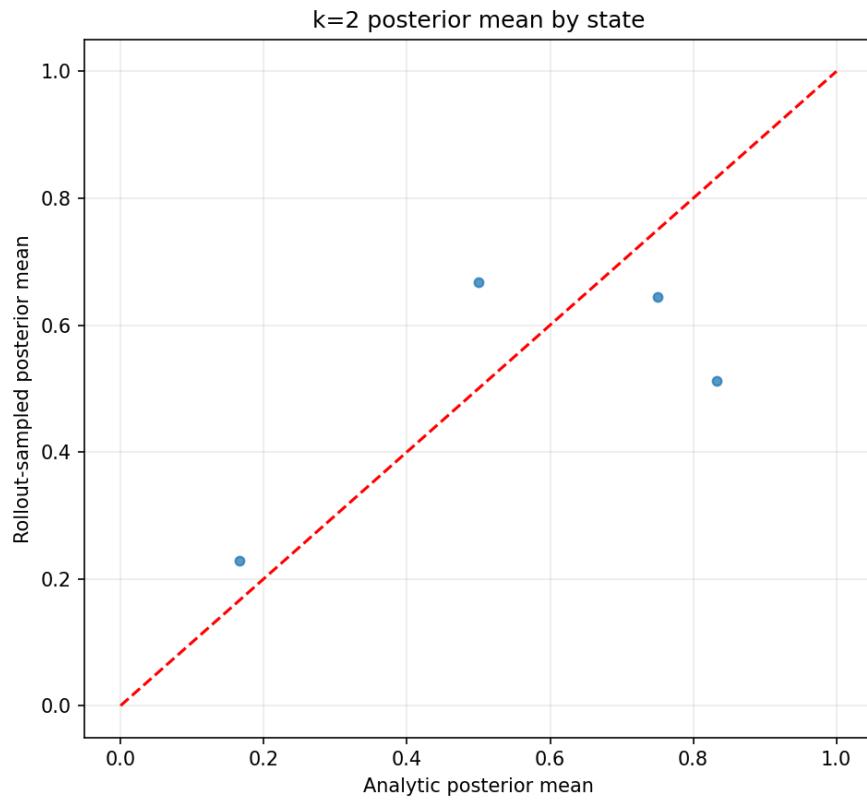


Figure 5: k=2: Step 2: posterior mean comparison

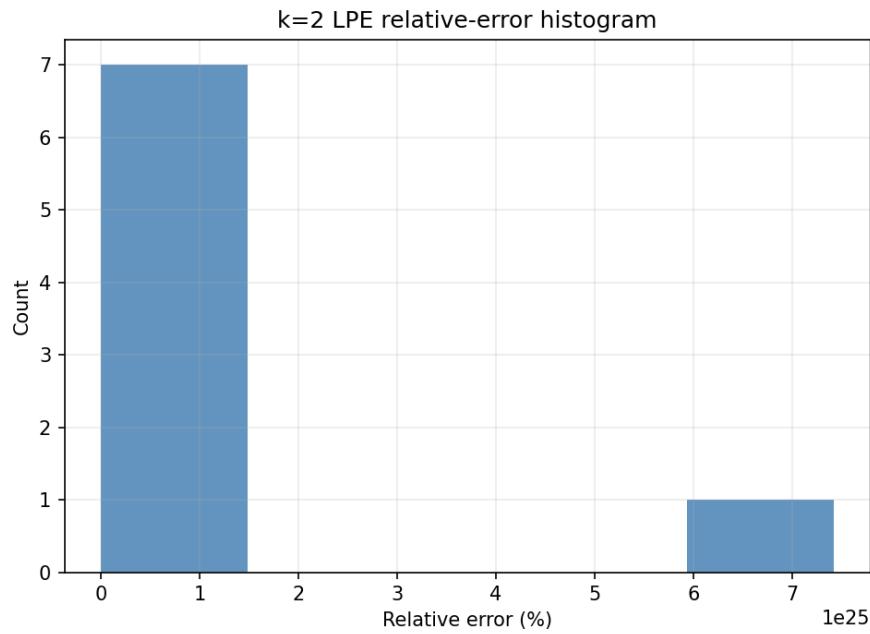


Figure 6: k=2: Step 3: LPE relative-error histogram

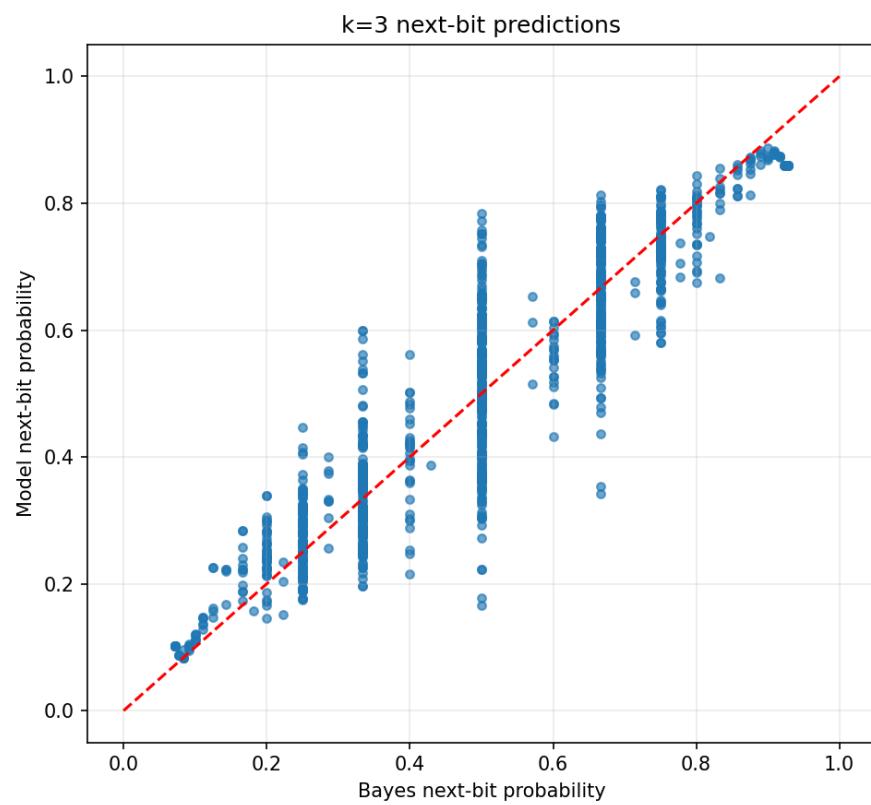


Figure 7: k=3: Step 1: next-bit prediction vs Bayes

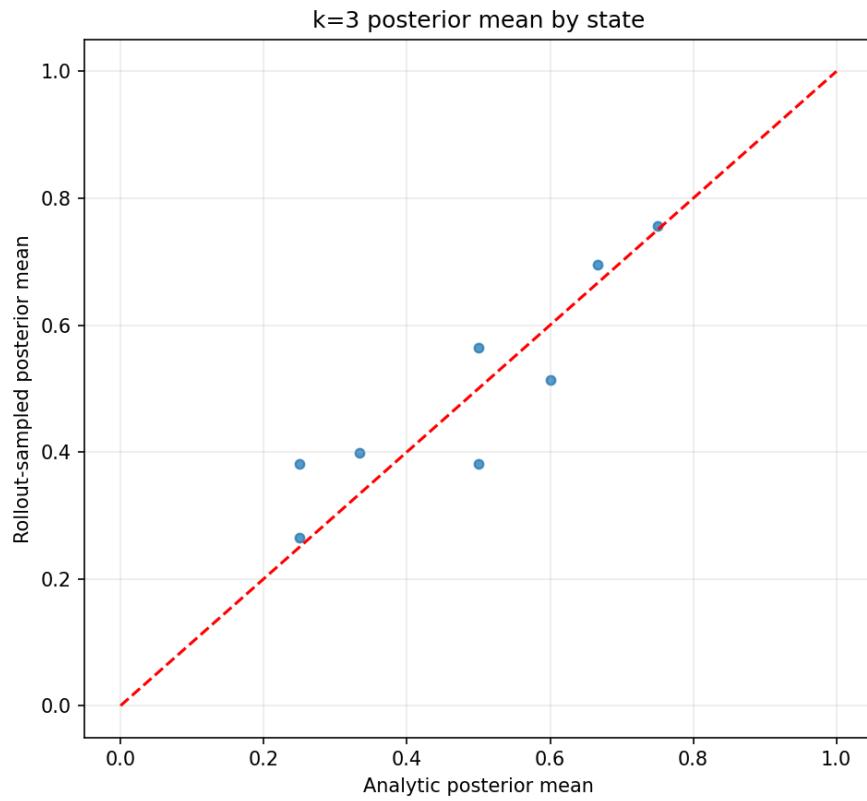


Figure 8: k=3: Step 2: posterior mean comparison

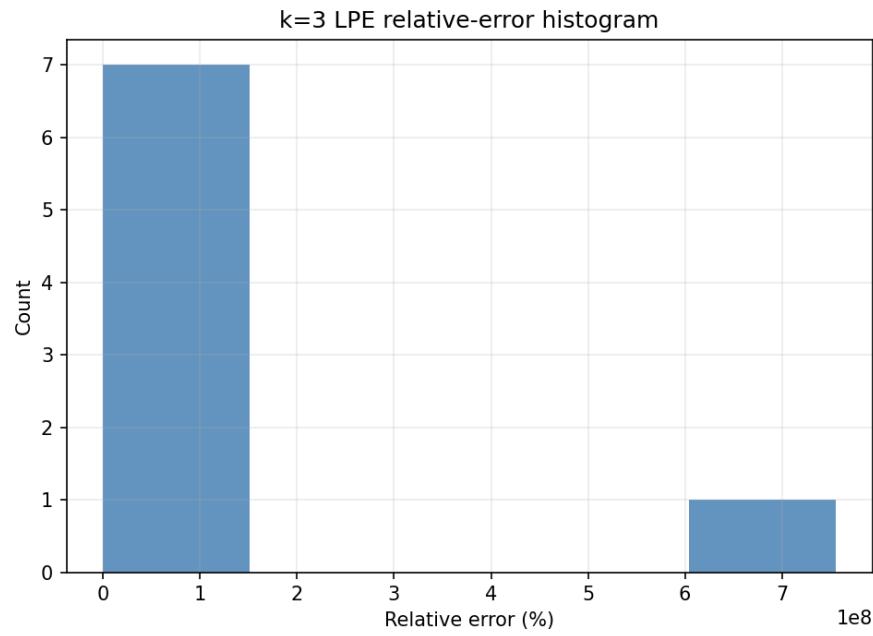


Figure 9: k=3: Step 3: LPE relative-error histogram